

walmart-busscase

April 4, 2024

Problem Statment : Walmart is an American multinational retail corporation that operates a chain of supercenters, discount departmental stores, and grocery stores in the United States. Walmart has more than 100 million customers worldwide. Wants to analyze the customer purchase behavior against the customer's gender and the various other factors to help the business make better decisions. Want to understand if the spending habits differ between male and female customers

```
[185]: #Import the dataset and do usual data analysis
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv("F:\\buss_cass\\data\\walmart_data.txt")
```

Checking the structure & characteristics of the dataset

```
[186]: df.head()
```

```
[186]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
0	1000001	P00069042	F	0-17	10	A	
1	1000001	P00248942	F	0-17	10	A	
2	1000001	P00087842	F	0-17	10	A	
3	1000001	P00085442	F	0-17	10	A	
4	1000002	P00285442	M	55+	16	C	

	Stay_In_Current_City_Years	Marital_Status	Product_Category	Purchase
0	2	0	3	8370
1	2	0	1	15200
2	2	0	12	1422
3	2	0	12	1057
4	4+	0	8	7969

```
[187]: df.shape
```

```
[187]: (550068, 10)
```

```
[188]: df.isna().sum()
```

```
[188]: User_ID          0
      Product_ID       0
      Gender           0
      Age              0
      Occupation       0
      City_Category    0
      Stay_In_Current_City_Years  0
      Marital_Status   0
      Product_Category  0
      Purchase         0
      dtype: int64
```

```
[189]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                           550068 non-null  object
2   Gender                               550068 non-null  object
3   Age                                   550068 non-null  object
4   Occupation                           550068 non-null  int64
5   City_Category                        550068 non-null  object
6   Stay_In_Current_City_Years          550068 non-null  object
7   Marital_Status                      550068 non-null  int64
8   Product_Category                    550068 non-null  int64
9   Purchase                            550068 non-null  int64
dtypes: int64(5), object(5)
memory usage: 42.0+ MB
```

```
[190]: #there are some int datatypes coverting into object
      dt = df.copy()
```

```
[191]: columns=['User_ID', 'Occupation', 'Marital_Status', 'Product_Category']
      dt[columns]=dt[columns].astype('object')
```

```
[192]: dt.describe(include = 'all')
```

```
[192]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
count	550068.0	550068	550068	550068	550068.0	550068	
unique	5891.0	3631	2	7	21.0	3	
top	1001680.0	P00265242	M	26-35	4.0	B	
freq	1026.0	1880	414259	219587	72308.0	231173	
mean	NaN	NaN	NaN	NaN	NaN	NaN	
std	NaN	NaN	NaN	NaN	NaN	NaN	

min	NaN	NaN	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN	NaN	NaN

	Stay_In_Current_City_Years	Marital_Status	Product_Category \
count	550068	550068.0	550068.0
unique	5	2.0	20.0
top	1	0.0	5.0
freq	193821	324731.0	150933.0
mean	NaN	NaN	NaN
std	NaN	NaN	NaN
min	NaN	NaN	NaN
25%	NaN	NaN	NaN
50%	NaN	NaN	NaN
75%	NaN	NaN	NaN
max	NaN	NaN	NaN

	Purchase
count	550068.000000
unique	NaN
top	NaN
freq	NaN
mean	9263.968713
std	5023.065394
min	12.000000
25%	5823.000000
50%	8047.000000
75%	12054.000000
max	23961.000000

```
[ ]:
```

```
[193]: df.nunique()
```

```
[193]: User_ID          5891
Product_ID         3631
Gender              2
Age                7
Occupation         21
City_Category      3
Stay_In_Current_City_Years  5
Marital_Status     2
Product_Category   20
Purchase          18105
dtype: int64
```

0.0.1 Observation:

1. In the get dataset there is no null values
2. There are total 5891 unique user and 3631 unique product_id
3. Most purchased product_id P00265242
4. UserId 1001680 has most visted customer
5. There are 7 unique age groups and most of the purchase belongs to age 26-35 group


Missing Value & Outlier Detection

```
[194]: df['Purchase'].describe()
```

```
[194]: count      550068.000000
      mean        9263.968713
      std         5023.065394
      min          12.000000
      25%         5823.000000
      50%         8047.000000
      75%        12054.000000
      max        23961.000000
      Name: Purchase, dtype: float64
```

The mean value 9263.96 and min and max values are 12.00 and 23961 as we know outlier will impact the mean value here mean value is very less compare to the max value

```
[195]: #Boxplot to find outliers
      plt.figure(figsize=(6, 4))
      sns.boxplot(data=df, x='Purchase', orient='h')
      plt.show()
```



```
[196]: #Removing the Outliers
def remove_outliers(data):
    # Calculate the 5% & 75%
    q5 = data.quantile(0.25)
    q95 = data.quantile(0.75)

    # Calculate the interquartile range
    iqr = q95 - q5

    # Define the lower and upper bounds for outliers
    lower_bound = q5 - 1.5 * iqr
    upper_bound = q95 + 1.5 * iqr

    clipped_data = np.clip(data, lower_bound, upper_bound)

    return clipped_data

# Example usage:
data = df['Purchase']
clipped_data = remove_outliers(data)
print("number of outliers: " + str(len(clipped_data)))
print("max outlier value: " + str(clipped_data.max()))
print("min outlier value: " + str(clipped_data.min()))
```

```
number of outliers: 550068
max outlier value:21400.5
min outlier value: 12.0
```

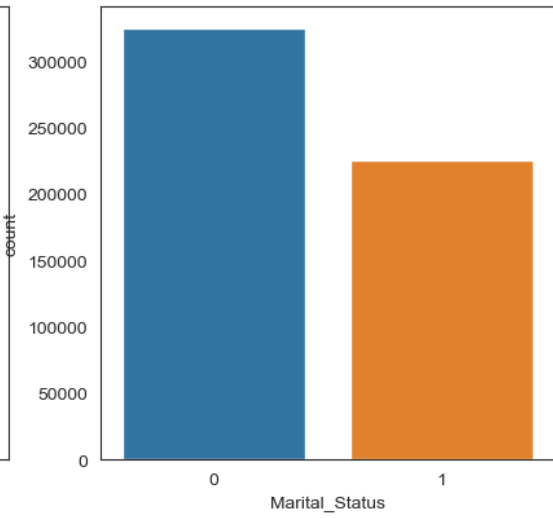
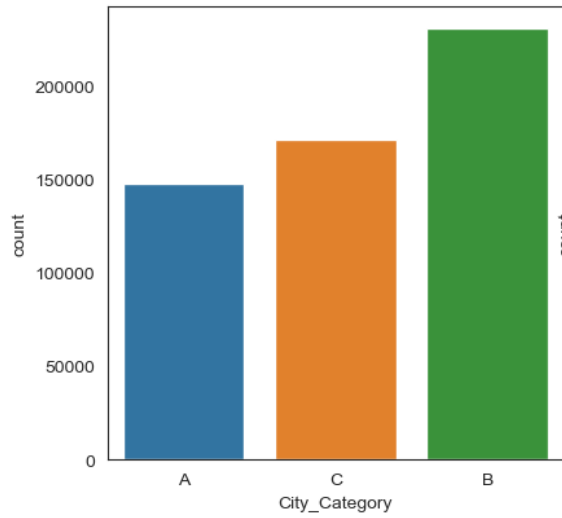


```
[ ]:
```


Data exploration

```
[ ]:
```

```
[ ]:
```

```
[197]: categorical_cols = ['Gender', 'Occupation', 'City_Category', 'Marital_Status', 'Product_Category']
fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(10, 10))
sns.countplot(data=df, x='Gender', ax=axs[0,0])
sns.countplot(data=df, x='Occupation', ax=axs[0,1])
sns.countplot(data=df, x='City_Category', ax=axs[1,0])
sns.countplot(data=df, x='Marital_Status', ax=axs[1,1])
plt.show()
plt.figure(figsize=(10, 10))
sns.countplot(data=df, x='Product_Category')
plt.show()
```





```
[198]: df_prod = df[['Product_Category' , 'Age']].value_counts().reset_index()
df_prod.head()
```

```
[198]:
```

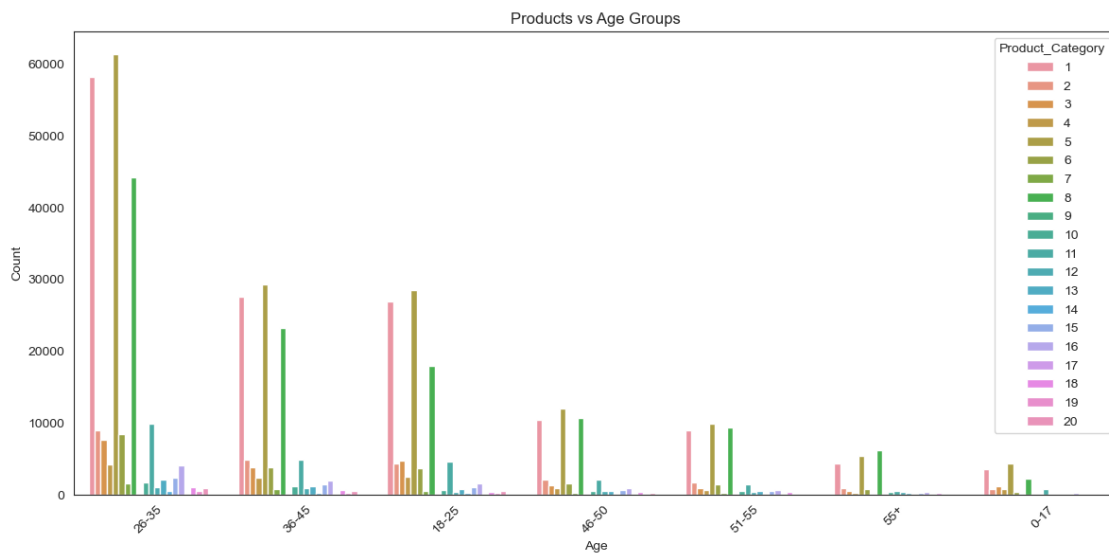
	Product_Category	Age	count
0	5	26-35	61473
1	1	26-35	58249
2	8	26-35	44256
3	5	36-45	29377
4	5	18-25	28522

```
[ ]:
```

```
[199]: #Visual Analysis-products and age groups
plt.figure(figsize=(12, 6))
```




```
sns.barplot(data=df_prod, x='Age', y='count', hue='Product_Category')
plt.title('Products vs Age Groups')
plt.xlabel('Age')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

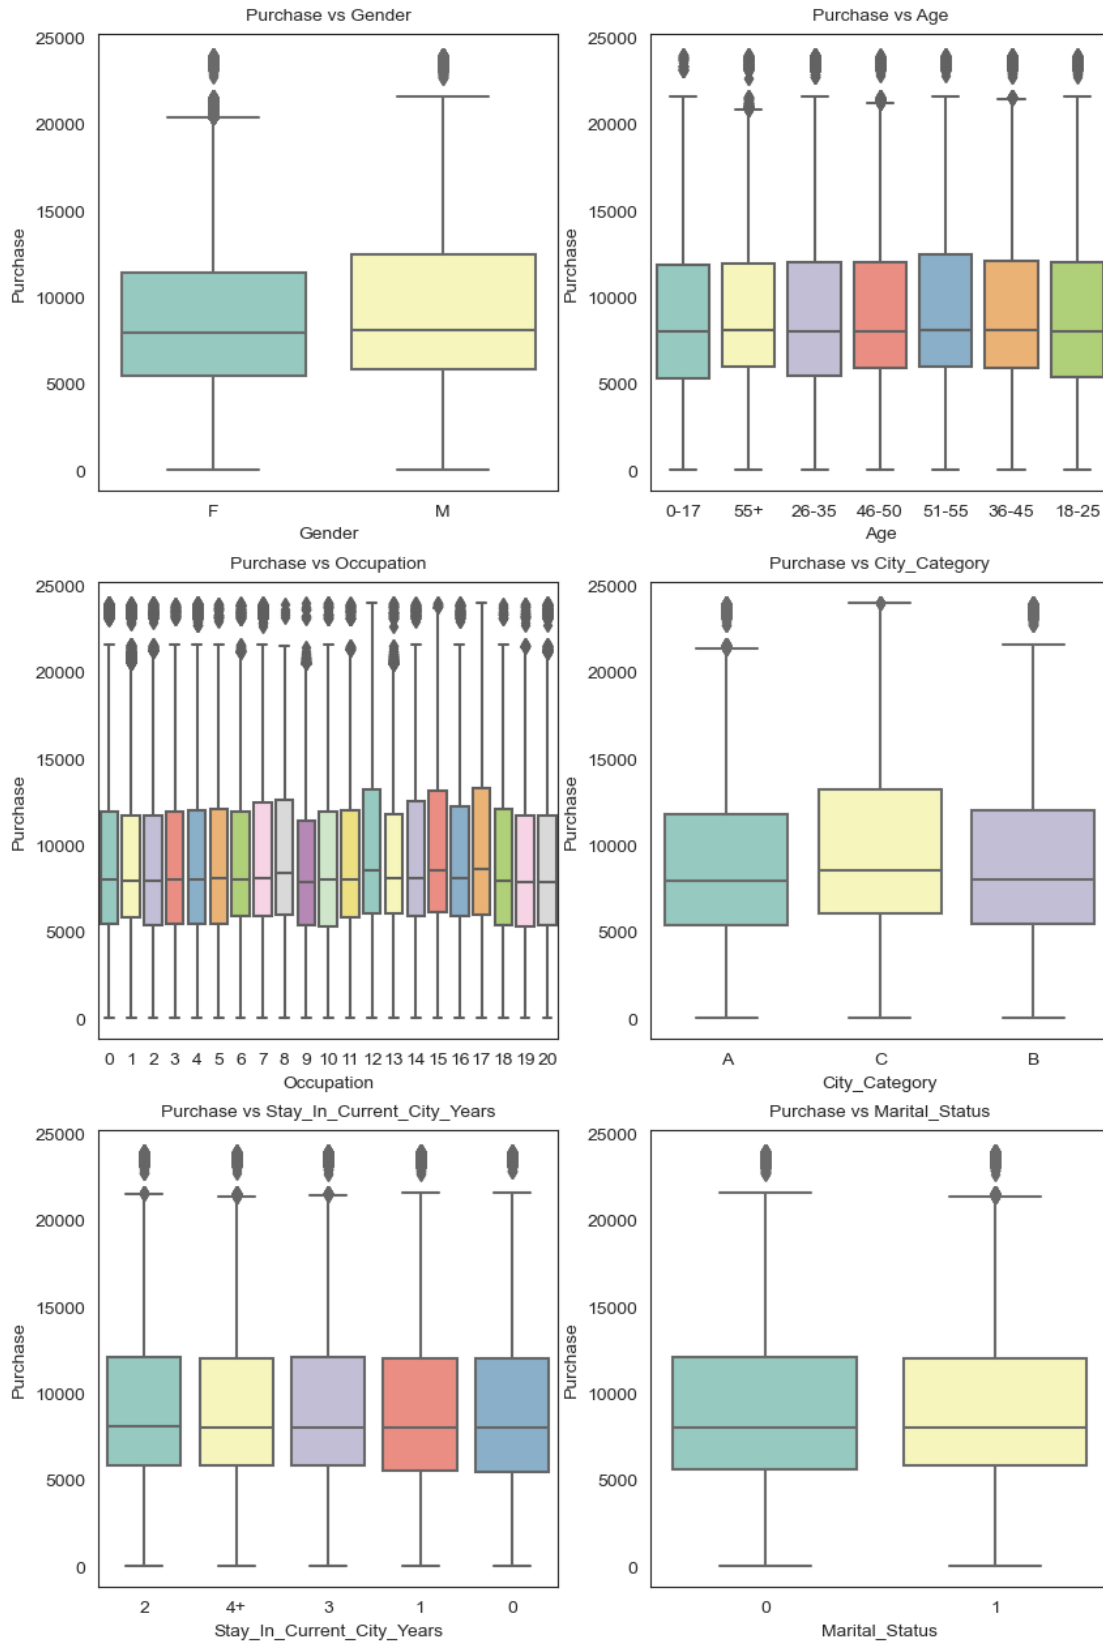


[200]: *#multivariate analysis between age,marital status, and the amount spent*


```
sns.boxplot(data=df, y='Purchase', x='Age', hue='Marital_Status')
plt.show()
```



```
[201]: #Visual Analysis
li = ['Gender', 'Age', 'Occupation', 'City_Category',
      'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category']
sns.set_style("white")
fig, axs = plt.subplots(nrows=3, ncols=2, figsize=(10,10))
fig.subplots_adjust(top=1.3)
count = 0
for row in range(3):
    for col in range(2):
        sns.boxplot(data=df, y='Purchase', x=li[count], ax=axs[row,col],
                    palette='Set3')
        axs[row,col].set_title(f"Purchase vs {li[count]}", pad=8, fontsize=10)
        count += 1
plt.show()
```



```
[202]: plt.figure(figsize=(10, 8))
sns.boxplot(data=df, y='Purchase', x=attrs[-1], palette='Set3')
plt.show()
```



Observations

1. Most of the users are Male
2. There are 20 different types of Occupation and Product_Category
3. More users belong to B City_Category
4. More users are Single as compare to Married
5. Product_Category - 1, 5, 8, & 11 have highest purchasing frequency.

Average amount spends per customer for Male and Female

```
[204]: cus_df = df.groupby(['User_ID', 'Gender'])[['Purchase']].sum().reset_index()
cus_df.head()
```

```
[204]:   User_ID Gender  Purchase
0  1000001      F    334093
```

1	1000002	M	810472
2	1000003	M	341635
3	1000004	M	206468
4	1000005	M	821001

```
[205]: # Gender wise value counts in avg_amt_df
cus_df['Gender'].value_counts()
```

```
[205]: Gender
M      4225
F      1666
Name: count, dtype: int64
```

```
[207]: male_avg = cus_df[cus_df['Gender']=='M']['Purchase'].mean()
female_avg = cus_df[cus_df['Gender']=='F']['Purchase'].mean()
print(male_avg)
print(female_avg)
```

925344.4023668639


712024.3949579832

mean purchase of Male = 925344.40, female = 712024.39, male spend more money than female

```
[227]: fig, axs = plt.subplots(nrows=1, ncols=2, figsize=(8,4))

# Plot histogram for 'Male'
sns.histplot(data=cus_df[cus_df['Gender'] == 'M'], x='Purchase', ax=axs[0],
             color='blue')
axs[0].set_title('Male')



# Plot histogram for 'Female'
sns.histplot(data=cus_df[cus_df['Gender'] == 'F'], x='Purchase', ax=axs[1],
             color='orange')
axs[1].set_title('Female')
plt.tight_layout()
plt.show()
```



```
[265]: male_df = cus_df[cus_df['Gender']=='M']
female_df = cus_df[cus_df['Gender']=='F']
genders = ["M", "F"]
male_sample_size = 1500
female_sample_size = 1000
male_means = []
female_means = []
for _ in range(1000):
    male_mean = male_df.sample(male_sample_size,replace=True)['Purchase'].mean()
    female_mean = female_df.sample(female_sample_size,replace=True)['Purchase'].
    ↪mean()

    male_means.append(male_mean)
    female_means.append(female_mean)
```

```
[266]: fig, axis = plt.subplots(nrows=1, ncols=2, figsize=(20, 6))
sns.histplot(male_means, kde=True, bins=35, ax=axis[0], color='skyblue')
sns.histplot(female_means, kde=True, bins=35, ax=axis[1], color='orange')
axis[0].set_title("Male - Distribution of means, Sample size: 1500")
axis[1].set_title("Female - Distribution of means, Sample size: 1000")
plt.show()
```



[]:

```
[276]: #Taking the values for z at 90%, 95% confidence interval
from scipy.stats import norm
z90 = norm.ppf(0.90)
z95 = norm.ppf(0.95)
```

Calculating 90% confidence interval for sample size male = 1500 & female = 1000:

```
[279]: #male
sample_mean_M = np.mean(male_means)
sample_std_M = pd.Series(male_means).std()

se_M = (z90*sample_std_M)/(np.sqrt(1500))

lower_limit = sample_mean_M - se_M
upper_limit = sample_mean_M + se_M
print("Male_CI_90:", [lower_limit, upper_limit] )
```

Male_CI_90: [924020.0866301863, 925703.0202684805]

```
[280]: #Female
sample_mean_F = np.mean(female_means)
sample_std_F = pd.Series(female_means).std()

se_F = (z90*sample_std_F)/(np.sqrt(1000))

lower_limit = sample_mean_F - se_F
upper_limit = sample_mean_F + se_F

print("Female_CI_90:", [lower_limit, upper_limit])
```

Female_CI_90: [712801.6347753559, 714966.6117366441]

Calculating 95% confidence interval for sample size male = 1500 & female = 1000:

```
[282]: #male
sample_mean_M = np.mean(male_means)
sample_std_M = pd.Series(male_means).std()

se_M = (z95*sample_std_M)/(np.sqrt(1500))

lower_limit = sample_mean_M - se_M
upper_limit = sample_mean_M + se_M
print("Male_CI_95:", [lower_limit, upper_limit] )
```

Male_CI_95: [923781.5424793141, 925941.5644193527]

```
[283]: #Female
sample_mean_F = np.mean(female_means)
sample_std_F = pd.Series(female_means).std()

se_F = (z95*sample_std_F)/(np.sqrt(1000))

lower_limit = sample_mean_F - se_F
upper_limit = sample_mean_F + se_F

print("Female_CI_95:", [lower_limit, upper_limit])
```

Female_CI_95: [712494.7643324954, 715273.4821795046]

CLT and Confidence interval considering marital status:

```
[288]: avg_Marital = df.groupby(['User_ID', 'Marital_Status'])[['Purchase']].sum().
        ↪reset_index()
avg_Marital.head()
```

```
[288]:   User_ID  Marital_Status  Purchase
0  1000001                0    334093
1  1000002                0    810472
2  1000003                0    341635
3  1000004                1    206468
4  1000005                1    821001
```

```
[289]: avg_Marital['Marital_Status'].value_counts()
```

```
[289]: Marital_Status
0      3417
1      2474
Name: count, dtype: int64
```

```
[287]: avgamt_married = avg_Marital[avg_Marital['Marital_Status']==1]
avgamt_single = avg_Marital[avg_Marital['Marital_Status']==0]
```



```

sample_size = 1000
num_repitions = 1000
married_means = []
single_means = []

for i in range(num_repitions):
    avg_married = avg_Marital[avg_Marital['Marital_Status']==1].
    ↪sample(sample_size, replace=True)['Purchase'].mean()
    avg_single = avg_Marital[avg_Marital['Marital_Status']==0].
    ↪sample(sample_size, replace=True)['Purchase'].mean()

    married_means.append(avg_married)
    single_means.append(avg_single)

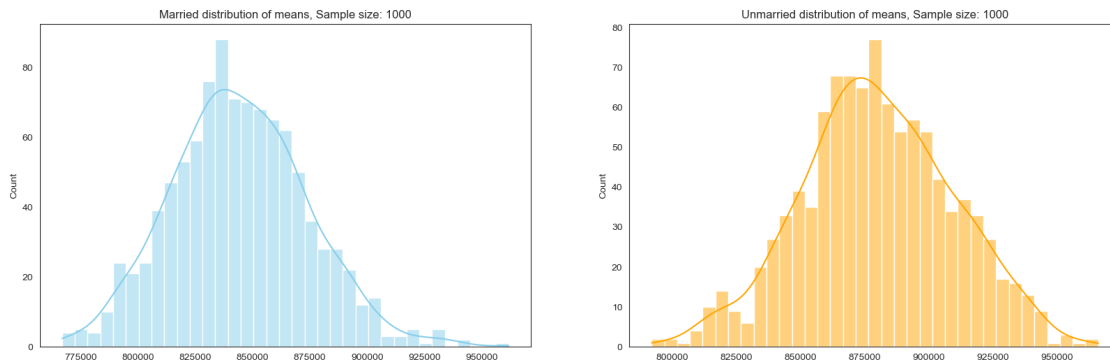
```

```

[285]: fig, axis = plt.subplots(nrows=1, ncols=2, figsize=(20, 6))
sns.histplot(married_means, kde=True, bins=35, ax=axis[0], color='skyblue')
sns.histplot(single_means, kde=True, bins=35, ax=axis[1], color='orange')
axis[0].set_title("Married distribution of means, Sample size: 1000")
axis[1].set_title("Unmarried distribution of means, Sample size: 1000")

plt.show()

```



Calculating 90% confidence interval for sample size married = 1000 & singles = 1000:

```

[291]: #married
sample_mean_M = np.mean(married_means)
sample_std_M = pd.Series(married_means).std()

se_M = (z90*sample_std_M)/(np.sqrt(1000))

lower_limit = sample_mean_M - se_M
upper_limit = sample_mean_M + se_M
print("Married_CI_90:", [lower_limit, upper_limit] )

```

Married_CI_90: [842975.4897113338, 845376.5821146662]

```
[292]: #single
sample_mean_s = np.mean(single_means)
sample_std_s = pd.Series(single_means).std()

se_s = (z90*sample_std_s)/(np.sqrt(1000))

lower_limit = sample_mean_s - se_s
upper_limit = sample_mean_s + se_s
print("Single_CI_90:", [lower_limit, upper_limit] )
```

Single_CI_90: [877405.7785519884, 879883.8127680115]

Calculating 95% confidence interval for sample size married = 1000 & singles = 1000:

```
[294]: #married
sample_mean_M = np.mean(married_means)
sample_std_M = pd.Series(married_means).std()

se_M = (z95*sample_std_M)/(np.sqrt(1000))

lower_limit = sample_mean_M - se_M
upper_limit = sample_mean_M + se_M
print("Married_CI_95:", [lower_limit, upper_limit] )
```

Married_CI_95: [842635.151545, 845716.920281]

```
[295]: #single
sample_mean_s = np.mean(single_means)
sample_std_s = pd.Series(single_means).std()

se_s = (z95*sample_std_s)/(np.sqrt(1000))

lower_limit = sample_mean_s - se_s
upper_limit = sample_mean_s + se_s
print("Single_CI_95:", [lower_limit, upper_limit] )
```

Single_CI_95: [877054.534418267, 880235.0569017329]

CLT and Confidence interval considering Age:

```
[296]: avgamt_age = df.groupby(['User_ID', 'Age'])[['Purchase']].sum().reset_index()
avgamt_age.head()
```

```
[296]:   User_ID  Age  Purchase
0  1000001  0-17    334093
1  1000002   55+    810472
2  1000003  26-35    341635
```

3	1000004	46-50	206468
4	1000005	26-35	821001

```
[298]: avgamt_age['Age'].value_counts()
```

```
[298]: Age
26-35    2053
36-45    1167
18-25    1069
46-50     531
51-55     481
55+       372
0-17      218
Name: count, dtype: int64
```

```
[319]: sample_dict = {}

age_intervals = ['26-35', '36-45', '18-25', '46-50', '51-55', '55+', '0-17']
for i in age_intervals:
    sample_dict[i] = []

for i in age_intervals:
    for j in range(200):


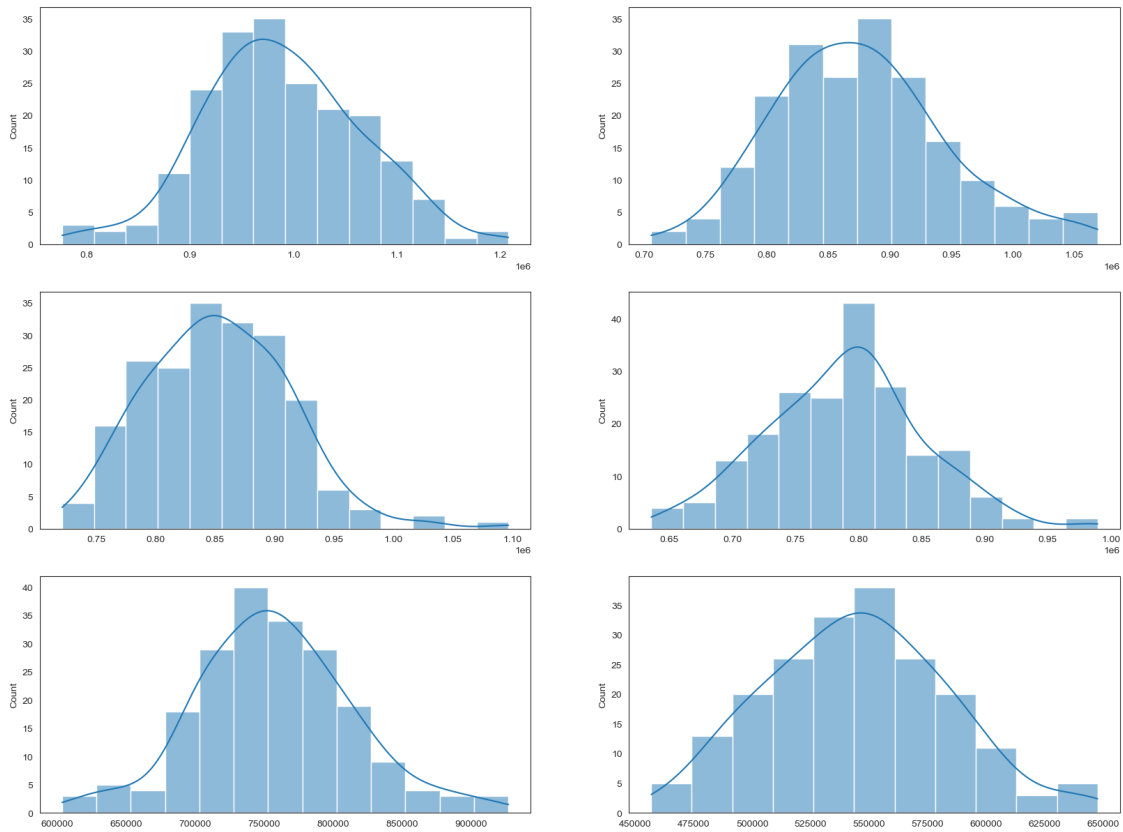
        mean = avgamt_age[avgamt_age['Age']==i].sample(sample_size,
↪replace=True)['Purchase'].mean()
        sample_dict[i].append(mean)
```

```
[320]: fig, axis = plt.subplots(nrows=3, ncols=2, figsize=(20, 15))

sns.histplot(sample_dict['26-35'],kde = True,ax=axis[0,0])
sns.histplot(sample_dict['36-45'],kde = True,ax=axis[0,1])
sns.histplot(sample_dict['18-25'],kde = True,ax=axis[1,0])
sns.histplot(sample_dict['46-50'],kde = True,ax=axis[1,1])
sns.histplot(sample_dict['51-55'],kde = True,ax=axis[2,0])
sns.histplot(sample_dict['55+'],kde = True,ax=axis[2,1])

plt.show()

plt.figure(figsize=(5, 3))
sns.histplot(sample_dict['0-17'],kde = True)
plt.show()
```



Calculating 90% confidence interval

```

[324]: all_population_means={}
all_sample_means = {}

sample_size = 200
num_repitions = 1000

age_intervals = ['26-35', '36-45', '18-25', '46-50', '51-55', '55+', '0-17']
for i in age_intervals:
    all_sample_means[i] = []
    all_population_means[i]=[]
    population_mean=avgamt_age[avgamt_age['Age']==i]['Purchase'].mean()
    all_population_means[i].append(population_mean)

for i in age_intervals:
    for j in range(num_repitions):

        mean = avgamt_age[avgamt_age['Age']==i].sample(sample_size,
↪replace=True)['Purchase'].mean()
        all_sample_means[i].append(mean)

for val in ['26-35', '36-45', '18-25', '46-50', '51-55', '55+', '0-17']:

    new_df = avgamt_age[avgamt_age['Age']==val]

    std_error = z90*new_df['Purchase'].std()/np.sqrt(len(new_df))
    sample_mean = new_df['Purchase'].mean()
    lower_lim = sample_mean - std_error
    upper_lim = sample_mean + std_error

    print("For age {} confidence interval of means: {:.2f}, {:.2f}").
↪format(val, lower_lim, upper_lim)

```

```

For age 26-35 confidence interval of means: (960481.20, 1018837.43)
For age 36-45 confidence interval of means: (842842.09, 916489.33)
For age 18-25 confidence interval of means: (820058.31, 889667.93)
For age 46-50 confidence interval of means: (740866.20, 844231.37)
For age 51-55 confidence interval of means: (716902.59, 809499.26)
For age 55+ confidence interval of means: (498668.64, 580725.85)
For age 0-17 confidence interval of means: (559232.93, 678502.69)

```

Calculating 95% confidence interval

```

[326]: all_population_means={}
all_sample_means = {}

sample_size = 200
num_repitions = 1000

```

```

age_intervals = ['26-35', '36-45', '18-25', '46-50', '51-55', '55+', '0-17']
for i in age_intervals:
    all_sample_means[i] = []
    all_population_means[i]=[]
    population_mean=avgamt_age[avgamt_age['Age']==i]['Purchase'].mean()
    all_population_means[i].append(population_mean)

for i in age_intervals:
    for j in range(num_repitions):

        mean = avgamt_age[avgamt_age['Age']==i].sample(sample_size,
↪replace=True)['Purchase'].mean()
        all_sample_means[i].append(mean)

for val in ['26-35', '36-45', '18-25', '46-50', '51-55', '55+', '0-17']:

    new_df = avgamt_age[avgamt_age['Age']==val]

    std_error = z95*new_df['Purchase'].std()/np.sqrt(len(new_df))
    sample_mean = new_df['Purchase'].mean()
    lower_lim = sample_mean - std_error
    upper_lim = sample_mean + std_error

    print("For age {} confidence interval of means: {:.2f}, {:.2f}").
↪format(val, lower_lim, upper_lim))

```

For age 26-35 confidence interval of means: (952209.61, 1027109.02)
 For age 36-45 confidence interval of means: (832403.10, 926928.32)
 For age 18-25 confidence interval of means: (810191.63, 899534.61)
 For age 46-50 confidence interval of means: (726214.90, 858882.66)
 For age 51-55 confidence interval of means: (703777.65, 822624.20)
 For age 55+ confidence interval of means: (487037.60, 592356.89)
 For age 0-17 confidence interval of means: (542327.27, 695408.35)

Recommendations:

1. Men spent more money than women, company can focus on retaining the male customers and getting more women.
2. Product_Category - 1, 5, 8 have highest purchasing frequency. it means these are the products which are most popular.
3. Unmarried customers spend more money than married customers, So company should focus on married customers.
4. Customers in the age 26-35 spend more money than the others, So company should focus on acquiring more customers in this age group.
5. We have more customers aged 26-35 in the city category B and A, company can focus more on this city category.

6. Some of the Product category like 19,20,13 have very less purchase. Company can think of dr
7. The top 10 users who have purchased more company should give more offers and discounts so th
8. The occupation which are contributing more company can think of offering credit cards or ot
9. The top products should be given focus in order to maintain the quality in order to further
10. People who are staying in city for an year have contributed to 35% of the total purchase an

[]:

[]: